

PROJECT CONTEXT DOCUMENT

Project Sentinel – Master Context (Version 1.0)

Purpose of this Document

This document contains the complete context, decisions, objectives, and current status of Project Sentinel. It should be provided at the beginning of every new ChatGPT conversation related to this project so that work can continue without losing context.

1. Project Objective

The objective is **not** to build another chatbot, recommendation engine, or generic AI application.

The objective is to build a **high-quality AI/Data Science project** that:

- Demonstrates strong AI and Machine Learning knowledge.
- Solves a genuine real-world problem.
- Is technically impressive to recruiters.
- Has research depth.
- Has startup potential.
- Can be completed by a single undergraduate student within approximately 4–6 months.
- Can be showcased during placements, internships, hackathons, and potentially serve as the foundation of a research paper.

The project should prioritize **technical quality and originality** over flashy features.

2. Ideation Journey

The project idea was not chosen randomly.

The following process was followed:

- Generated hundreds of problems across multiple industries using Gemini, Claude, Perplexity, DeepSeek, Kimi AI, and ChatGPT.
- Focused only on problems that could be solved using AI and Data Science (no hardware).
- Eliminated generic domains and over-explored student application ideas.
- Reduced the list through multiple elimination rounds based on:
 - Novelty
 - Real-world impact
 - AI depth

- Data availability
- Evaluation feasibility
- Recruiter appeal
- Compared the strongest candidates.
- Explored AI Reliability and AI Observability as a new-generation AI domain.
- Conducted multiple validation rounds.

Final shortlisted direction:

AI Reliability / AI Infrastructure / AI Observability

3. Final Working Project

Project Name

Project Sentinel

(Currently the working name.)

Working Title

Semantic Reliability Inference Framework for Production Large Language Models

Elevator Pitch

Project Sentinel is an AI reliability framework that detects semantic performance degradation in deployed AI systems before humans notice, even when no ground-truth labels are available.

Instead of generating AI responses, Sentinel evaluates whether deployed AI systems are becoming less reliable over time.

4. Core Problem

Modern AI systems often fail silently.

Traditional monitoring systems can detect:

- CPU usage
- Memory

- Latency
- API failures

But they cannot reliably detect:

- Semantic degradation
- Declining reasoning quality
- Silent hallucination increase
- Confidence drift
- Behavioral degradation

The project attempts to solve this problem.

5. Vision

Current AI systems answer:

"What is the answer?"

Project Sentinel answers:

"Can this AI still be trusted?"

6. Current Hypothesis

Current AI observability systems perform monitoring using:

- Human evaluation
- Benchmark evaluation
- Rule-based metrics
- Logging
- Tracing

The hypothesis is that AI systems can detect semantic reliability degradation using multiple behavioral signals without requiring continuous ground-truth labels.

This hypothesis still requires validation through literature review.

7. Current Proposed Solution

The proposed system continuously monitors model outputs and extracts multiple behavioral signals.

Instead of relying on one metric, Sentinel combines evidence from several sources to infer semantic reliability.

Potential signals include:

- Semantic consistency
- Response stability
- Embedding-space drift
- Confidence behavior
- Calibration indicators
- Output entropy proxies
- Behavioral evolution over time

These signals are fused into a Semantic Reliability Score.

8. High-Level Workflow

User Query



AI Model



Sentinel Monitoring Layer



Feature Extraction



Semantic Reliability Inference Engine



Reliability Score



Early Warning Alerts

9. Key AI Domains

The project combines:

- Natural Language Processing
 - Machine Learning
 - Representation Learning
 - Embedding Analysis
 - Statistical Monitoring
 - Drift Detection
 - AI Evaluation
 - AI Observability
 - Reliability Engineering
-

10. Expected Datasets

Potential benchmark datasets include:

- TruthfulQA
- GSM8K
- MMLU
- HumanEval

Potential degradation scenarios include:

- Prompt corruption
 - Distribution shift
 - Context truncation
 - Quantization
 - Model updates
-

11. Evaluation Metrics

The project should be evaluated using measurable metrics.

Examples:

- Detection Accuracy
- Precision
- Recall
- False Positive Rate
- Detection Lead Time

- Robustness across degradation scenarios
-

12. Validation Completed

The following questions have already been analyzed.

Problem Validation

Result:

PASS

The problem is real and increasingly important.

Industry Need

PASS

LLM reliability is becoming a major production concern.

Recruiter Value

PASS

The project stands out significantly compared to common student projects.

Dataset Availability

PASS

Suitable benchmark datasets appear available.

Student Buildability

PASS

A meaningful MVP is considered feasible.

Existing Solutions

PARTIALLY VALIDATED

Existing observability platforms solve adjacent problems.

However, the exact research gap still needs careful literature validation.

Algorithm Validation

PARTIALLY VALIDATED

Behavioral fingerprinting already exists in literature.

The novelty should not be behavioral fingerprinting itself.

Instead, Sentinel should focus on semantic reliability inference rather than merely detecting behavioral changes.

13. Current Risks

The biggest risks are:

- Becoming only a monitoring dashboard.
- Reproducing existing observability tools.
- Lacking a clearly differentiated algorithmic contribution.
- Failing to define Semantic Reliability mathematically.

These risks should remain a primary focus during future discussions.

14. Design Philosophy

This project should NOT become:

- Another chatbot.
- Another RAG application.
- Another dashboard.
- Another LLM wrapper.

The project should focus on creating a meaningful AI reasoning or inference component.

The algorithm is more important than the interface.

15. Remaining Research Tasks

Remaining work includes:

1. Complete literature review.
2. Competitor analysis.
3. Final research gap definition.
4. Mathematical definition of Semantic Reliability.
5. Design of the Semantic Reliability Inference Engine.
6. Dataset strategy.
7. Experimental methodology.
8. System architecture.
9. MVP implementation.
10. Evaluation.

16. Collaboration Rules

When providing suggestions:

- Challenge assumptions instead of agreeing automatically.
- Prefer evidence over speculation.
- Reject weak ideas quickly.
- Prioritize originality, technical depth, and practical feasibility.
- Treat this project like a research project rather than a classroom assignment.
- Recommend pivots if stronger evidence emerges.

17. Long-Term Goal

The desired outcome is a project that:

- Impresses AI/ML recruiters.
- Demonstrates advanced AI engineering skills.
- Can be expanded into a startup or research publication.
- Becomes the strongest project in the portfolio.

The goal is not simply to finish a project, but to create a technically credible AI system with measurable impact.

Current Status

Phase: Research Validation

Overall Confidence: High

Recommendation: Continue with Project Sentinel while refining the research gap and algorithmic contribution through systematic literature review and iterative design.